

# Sanika Killekar

Boston, MA | (857) 294-8020 | [killekar.s@northeastern.edu](mailto:killekar.s@northeastern.edu) | [GitHub](#) | [LinkedIn](#)

## SUMMARY

I build end to end data engineering and AI systems: I design reliable data pipelines in Python and SQL with Airflow, Spark, Kafka, and cloud storage, and I also build agentic AI and RAG applications that help teams search, summarize, and act on information with grounded outputs. I focus on production readiness through testing, CI/CD, monitoring, and secure deployments.

## WORK EXPERIENCE

### InfinityPool Finnotech Pvt Ltd - Machine Learning Co-op May 2024 – May 2025

- Led an end-to-end agentic RAG pipeline in **LangGraph** (routing, tool-calls, retries) to generate grounded next-day market summaries; processed **48k+ docs/day** and improved signal accuracy by **15%**.
- Served the LLM using **vLLM** on GPU and reduced GPU utilization by **28%** by eliminating redundant requests.
- Orchestrated ingestion with **Apache Airflow and Docker**; streamed feeds via **Kafka**; stored raw docs in **S3** and structured data in **PostgreSQL**, achieved **99.2%** successful DAG runs
- Built an internal analyst copilot using RAG with **ChromaDB** and tool-calling (SQL and document search) plus guardrails;
- supported **50+ daily users** and reduced repeat queries by **27%**.
- Deployed on **Kubernetes** and monitored with **Prometheus and Grafana**; tracked ingestion lag, RAG latency, and uptime.

### Tata Institute of Fundamental Research (TIFR) - AI Research Intern

*May 2023 – Apr 2024*

- Deployed a citation-first research copilot (**LangChain , FAISS, and GPT-4-turbo**) that searches, verifies, and summarizes evidence across **90K+ internal papers**, reducing literature review time by **35%**.
- Fine-tuned using **LoRA/QLoRA** to improve response consistency and reduce unsafe outputs; validated improvements with repeatable evaluation runs.
- Replaced **Azure Form Recognizer** with Unstructured.io for PDF parsing at scale, **lowering per-document processing cost**.
- Designed automated ETL pipelines (**Azure Data Factory and PySpark**) to ingest and clean heterogeneous sources (PDFs, patents, SQL DBs), reducing preprocessing time by **31%**.
- Containerized and deployed the RAG app on **AKS** using **Docker**, integrating authentication via **Azure Key Vault**.
- Automated CI/CD with **GitHub Actions** (unit tests), cutting deployment time from **45 to 10 mins**.

## EDUCATION

### Northeastern University, Boston, MA

*Aug 2025 – May 2027*

*Master of Science in Data Science*

**Relevant Coursework:** Machine Learning, MLOps, Algorithms, Database Management Systems, Data Engineering

### University of Mumbai, India

*Aug 2021 – May 2025*

*Bachelor of Engineering in Information Technology*

## TECHNICAL SKILLS

**Programming:** Python, SQL, Bash, R | **Operating Systems:** Linux, Windows | **Analytics:** Power BI, Tableau, Looker, Excel

**Machine Learning:** PyTorch, TensorFlow, Scikit-learn, XGBoost, Pandas, NumPy, NLP

**LLMs & RAG:** LangChain, LangGraph, OpenAI API, FAISS, ChromaDB, vLLM, FastAPI

**Data Engineering:** PySpark, Databricks, PostgreSQL, MySQL, ClickHouse, S3, Snowflake

**Cloud & MLOps:** AWS, Azure, GCP, Docker, Kubernetes, Apache Airflow, Kafka, GitHub Actions, Terraform, Prometheus, Grafana

## PROJECTS

### HubScout - Large-Scale Semantic Retrieval System

*Python, FastAPI, LangChain, ChromaDB/FAISS, vLLM, GCP*

- Built a semantic retrieval and RAG service over **350k+ records**; improved search relevance using embedding and reranking.
- Deployed the **FastAPI** service with GPU-backed inference (**vLLM**); added **agent behavioral regression tests** (tool-use accuracy, grounding rate, harmful-output checks) with red-team prompts to keep releases stable.

### IMDb ML Data Pipeline

*Azure Data Factory, Databricks (PySpark), Snowflake, dbt, Great Expectations*

- Built Azure Medallion (Bronze/Silver/Gold) pipeline to load **180M+ IMDb rows** into Snowflake, cutting prep time by **65%**.
- Added dbt and Great Expectations checks and shipped **Power BI** dashboards, improving data reliability by **85%**.

## PUBLICATIONS & INTELLECTUAL PROPERTY

### Predictive Analysis Based on Leading Parameters for Stock Market, IEEE ICNTE-2026

*Nov 2025*

### Crowd Flow Analyzer (Copyright Registered), Govt. of India

*Jul 2024*